

Challenge-2

Wong Zi Xin

2023-08-21

Welcomel Hope you have watched the lecture videos and followed the instructions in code-along. Go through the steps described below, carefully. It is totally fine to get stuck - **ASK FOR HELP**: reach out to your friends, TAs, or the discussion forum on Canvas.

Here is what you have to do,

1. Pair with a neighbor and work
2. Download the Challenge-2.Rmd and playlist_data.csv files from Canvas
3. Move the downloaded files to the folder, "Week-2"
4. Set it as the working directory
5. Edit content wherever indicated
6. Remember to set eval=TRUE after completing the code to generate the output
7. Ensure that echo=TRUE so that the code is rendered in the final document
8. Inform the tutor/instructor upon completion
9. Submit the document on Canvas after they approve
10. Attendance will be marked only after submission
11. Once again, **do not hesitate** to reach out to the tutors/instructor, if you are stuck

I. Exploring music preferences

A. Background

Imagine that you have been hired as a data analyst by a radio station to analyze music preferences of their DJs. They have provided you with a dataset, playlist_data.csv, containing information about DJs, their preferred music genres, song titles, and ratings.

Using the data-set you are required to complete some tasks that are listed subsequently. All these tasks are based on the concepts taught in the video lectures. The questions may not be entirely covered in the lectures; To complete them, you are encouraged to use Google and the resources therein.

B.Tasks

Task-1

In the lecture, we used two data-sets, starwars and anscombe's quartet that were readily available with the packages, tidyverse and taise, respectively. When we have to use custom-made data-sets or the ones like we downloaded from Canvas, we have to import it using the R commands before using them. All the questions below are related to this task.

Question 1.1: What does the term "CSV" in playlist_data.csv stand for, and why is it a popular format for storing tabular data?

Solution:
The word "CSV" stands for "Comma-Separated Values". It is popular because each line represents a record and it is easy to read and write for everyone.

Question 1.2: load the tidyverse package to work with .csv files in R.

Solution:

```
# Load the necessary package to work with CSV files in R.
library(tidyverse)

## -- Attaching core tidyverse packages -- tidyverse 2.0.0 --
## ✓ dplyr      1.1.2 ✓ readr      2.1.4
## ✓ forcats    1.0.0 ✓ stringr   1.5.0
## ✓ ggplot2    3.4.3 ✓ tibble    3.2.1
## ✓ lubridate  1.9.2 ✓ tidyr     1.3.0
## ✓ purrr      1.0.2
## --- Conflicts --- tidyverse_conflicts() ---
## * dplyr::filter() masks stats::filter()
## * dplyr::lag()   masks stats::lag()
## * Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

Question 1.3: Import the data-set, playlist_data.csv

Solution:

```
# Import the "playlist_data.csv" dataset into R
read_csv("playlist_data.csv")

## Rows: 26 Columns: 7
## -- Column specification --
## Delimiter: ","
## chr (4): DJ_Name, Music_Genre, Experience, Location
## dbl (3): Rating, Age, Plays_Per_Week
##
## I use 'spec()' to retrieve the full column specification for this data.
## I specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
## # A tibble: 26 × 7
##   DJ_Name Music_Genre Rating Experience      Age Location Plays_Per_Week
##   <chr>   <chr>       <dbl> <chr>      <dbl> <chr>      <dbl>
## 1 DJ A    Pop           4.2 Advanced  28 City X      80
## 2 DJ B    Rock          3.8 Intermediate 24 City Y      60
## 3 DJ C    Electronic    4.5 Advanced  30 City Z      100
## 4 DJ D    Pop           4 Intermediate 22 City X      70
## 5 DJ E    Electronic    4.8 Advanced  27 City Y      90
## 6 DJ F    Rock          3.6 Intermediate 25 City Z      55
## 7 DJ G    Pop           4.3 Advanced  29 City X      85
## 8 DJ H    Electronic    4.1 Intermediate 23 City Y      75
## 9 DJ I    Rock          3.9 Advanced  31 City Z      70
## 10 DJ J   Pop           4.4 Intermediate 26 City X      95
## # 16 more rows
```

Question 1.4: Assign the data-set to a variable, playlist_data

Solution:

```
# Assign the variable to a dataset
playlist_data <- read_csv("playlist_data.csv")

## Rows: 26 Columns: 7
## -- Column specification --
## Delimiter: ","
## chr (4): DJ_Name, Music_Genre, Experience, Location
## dbl (3): Rating, Age, Plays_Per_Week
##
## I use 'spec()' to retrieve the full column specification for this data.
## I specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

From now on, you can use the name of the variable to view the contents of the data-set

Question 1.5: Get more information about read_csv() command and provide a screenshot of the information displayed in the "Help" tab of the "Files" pane

Solution:

```
# More information about the R command, complete the code
?read_csv("playlist_data.csv")

knitr::include_graphics("Screenshot.png")

read_delim(readr) R Documentation
```

Read a delimited file (including CSV and TSV) into a tibble

Description

read_csv() and read_tsv() are special cases of the more general read_delim(). They're useful for reading the most common types of flat file data, comma separated values and tab separated values, respectively. read_csv2() uses ; for the field separator and , for the decimal point. This format is common in some European countries.

Usage

```
read_delim(
  file,
  delim = NULL,
  quote = "\"",
  escape_backslash = FALSE,
  escape_double = TRUE
)
```

More information about read_csv command

Question 1.6: What does the skip argument in the read_csv() function do?

Solution: It is used to specify the number of lines to skip at the beginning of the CSV file before reading in the data.

Question 1.7: Display the contents of the data-set

Solution:

```
# Type the name of the variable, to see what it contains
playlist_data

## # A tibble: 26 × 7
##   DJ_Name Music_Genre Rating Experience      Age Location Plays_Per_Week
##   <chr>   <chr>       <dbl> <chr>      <dbl> <chr>      <dbl>
## 1 DJ A    Pop           4.2 Advanced  28 City X      80
## 2 DJ B    Rock          3.8 Intermediate 24 City Y      60
## 3 DJ C    Electronic    4.5 Advanced  30 City Z      100
## 4 DJ D    Pop           4 Intermediate 22 City X      70
## 5 DJ E    Electronic    4.8 Advanced  27 City Y      90
## 6 DJ F    Rock          3.6 Intermediate 25 City Z      55
## 7 DJ G    Pop           4.3 Advanced  29 City X      85
## 8 DJ H    Electronic    4.1 Intermediate 23 City Y      75
## 9 DJ I    Rock          3.9 Advanced  31 City Z      70
## 10 DJ J   Pop           4.4 Intermediate 26 City X      95
## # 16 more rows
```

Question 1.8: Assume you have a CSV file named sales_data.csv containing information about sales transactions. How would you use the read_csv() function to import this file into R and store it in a variable named sales_data?

Solution:

```
# No output is required for this code
# Only the list of commands that execute the task mentioned in the question are required
library(tidyverse)
read_csv("sales_data.csv")
sales_data <- read_csv("sales_data.csv")
```

Task-2

After learning to import a data-set, let us explore the contents of the data-set through the following questions

Question 2.1: Display the first few rows of the data-set to get an overview of its structure

Solution:

```
# Type the name of the variable we assigned the data-set to
head(playlist_data)

## # A tibble: 6 × 7
##   DJ_Name Music_Genre Rating Experience      Age Location Plays_Per_Week
##   <chr>   <chr>       <dbl> <chr>      <dbl> <chr>      <dbl>
## 1 DJ A    Pop           4.2 Advanced  28 City X      80
## 2 DJ B    Rock          3.8 Intermediate 24 City Y      60
## 3 DJ C    Electronic    4.5 Advanced  30 City Z      100
## 4 DJ D    Pop           4 Intermediate 22 City X      70
## 5 DJ E    Electronic    4.8 Advanced  27 City Y      90
## 6 DJ F    Rock          3.6 Intermediate 25 City Z      55
```

Question 2.2: Display all the columns of the variable stacked one below another

Solution:

```
# Stack columns of playlist_data
glimpse(playlist_data)

## Rows: 26
## Columns: 7
## $ DJ_Name      <chr> "DJ A", "DJ B", "DJ C", "DJ D", "DJ E", "DJ F", "DJ G",...
## $ Music_Genre  <chr> "Pop", "Rock", "Electronic", "Pop", "Electronic", "Rock...
## $ Rating       <dbl> 4.2, 3.8, 4.5, 4.0, 4.8, 3.6, 4.3, 4.1, 3.9, 4.4, 4.6,...
## $ Experience   <chr> "Advanced", "Intermediate", "Advanced", "Intermediate",...
## $ Age         <dbl> 28, 24, 30, 22, 27, 25, 29, 23, 31, 26, 32, 28, 29, 25,...
## $ Location     <chr> "City X", "City Y", "City Z", "City Z", "City Y", "City...
## $ Plays_Per_Week <dbl> 80, 60, 100, 70, 90, 55, 85, 75, 70, 95, 110, 75, 60, 8...
```

Question 2.3: How many columns are there in the dataset?

Solution:

```
# Number of columns
ncol(playlist_data)

## [1] 7
```

Question 2.4: What is the total count of DJs?

Solution:

```
# Number of DJs
nrow(playlist_data)

## [1] 26
```

Question 2.5: Display all the location of all the DJs

Solution:

```
# Location of DJs
playlist_data$Location

## [1] "City X" "City Y" "City Z" "City X" "City Y" "City Z" "City X" "City Y"
## [9] "City Z" "City X" "City Y" "City Z" "City X" "City Y" "City X"
## [17] "City Y" "City Z" "City X" "City Y" "City Z" "City X" "City Y" "City Z"
## [25] "City X" "City Y"
```

Question 2.6: Display the age of the DJs

Solution:

```
# Age of DJs
playlist_data$Age

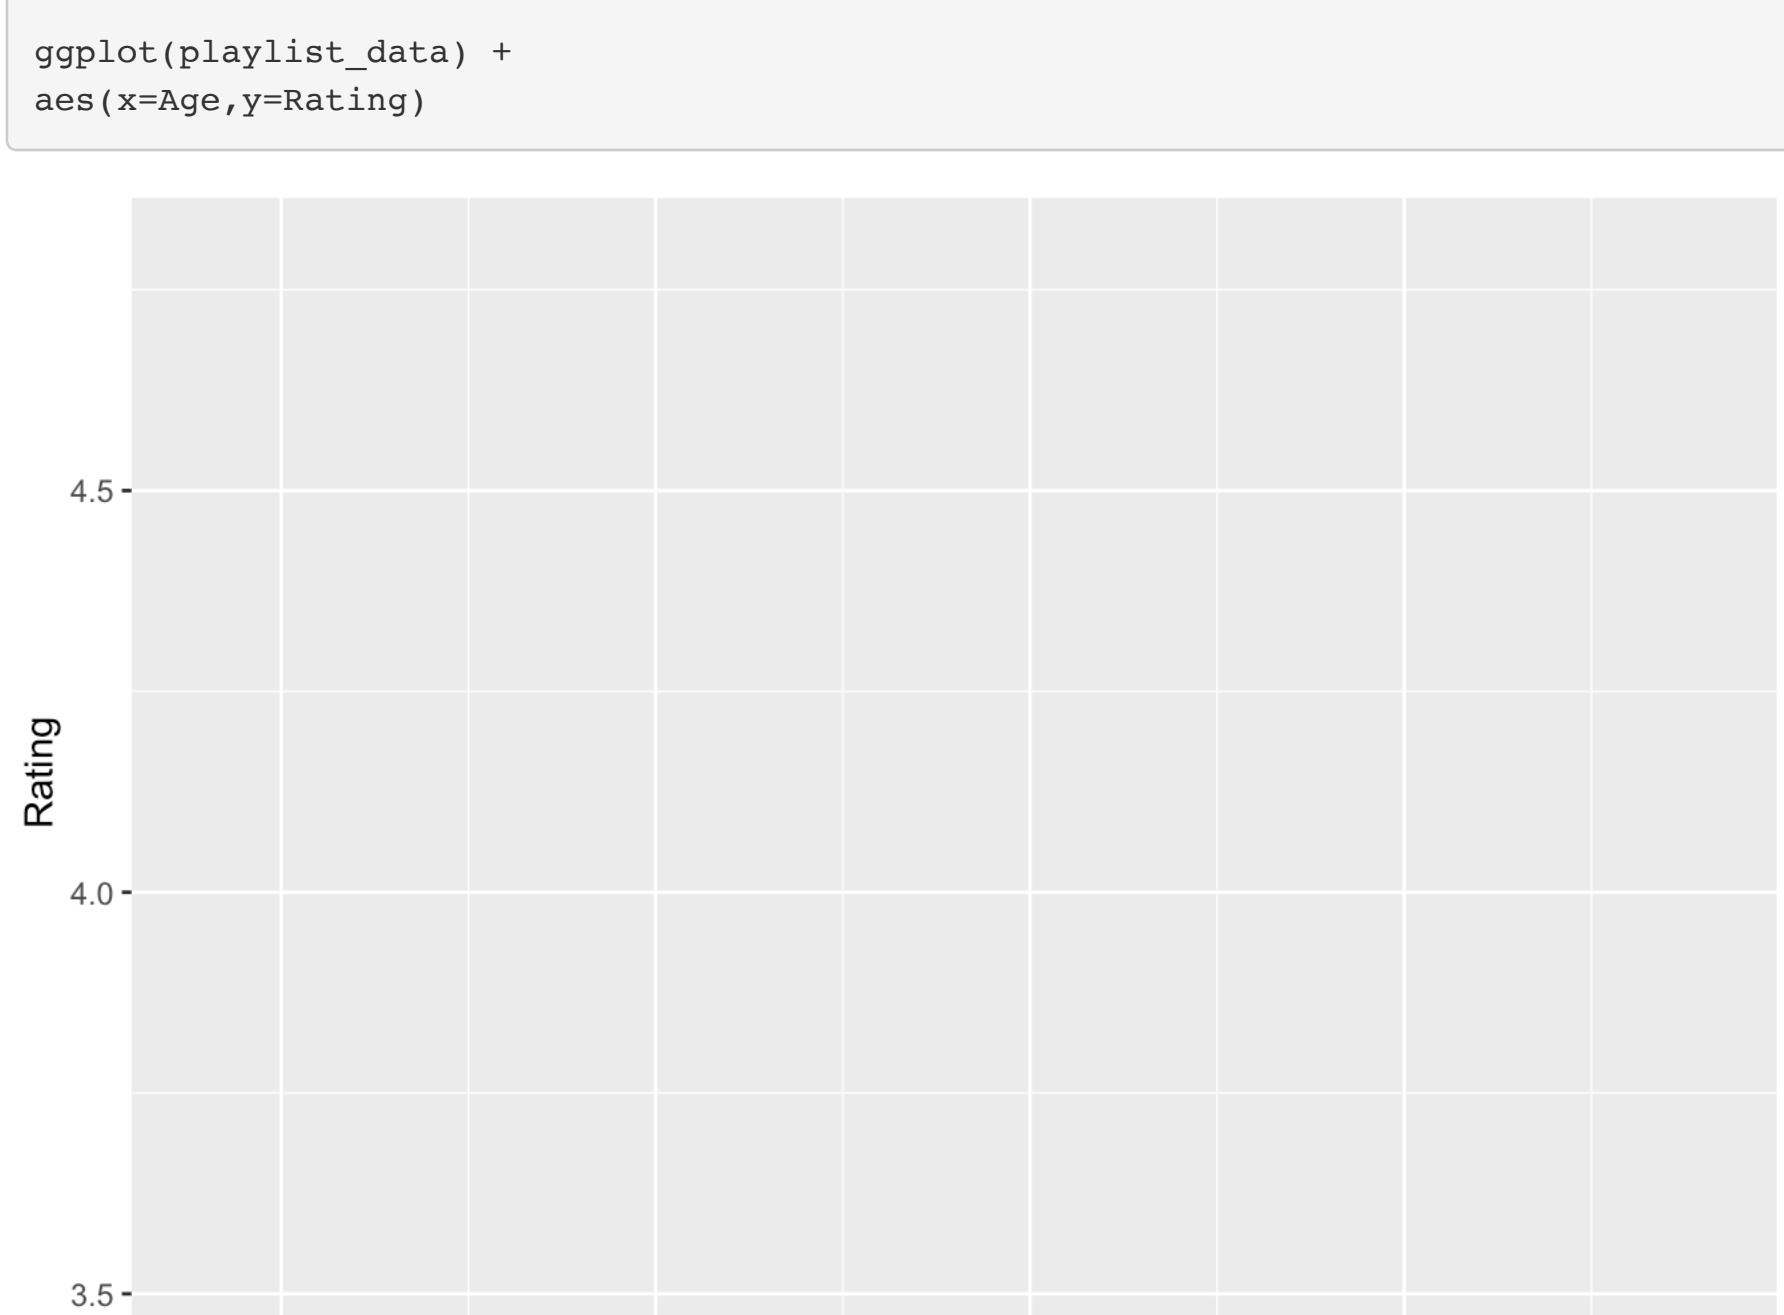
## [1] 28 24 30 22 27 25 29 23 31 26 32 28 29 25 31 26 27 24 29 23 28 24 30 22 27
## [26] 25
```

Task-3

Let us plot the data to get more insights about the DJs.

Question 3.1: Create a plot to visualize the relationship between DJs' ages and their ratings.

Solution:



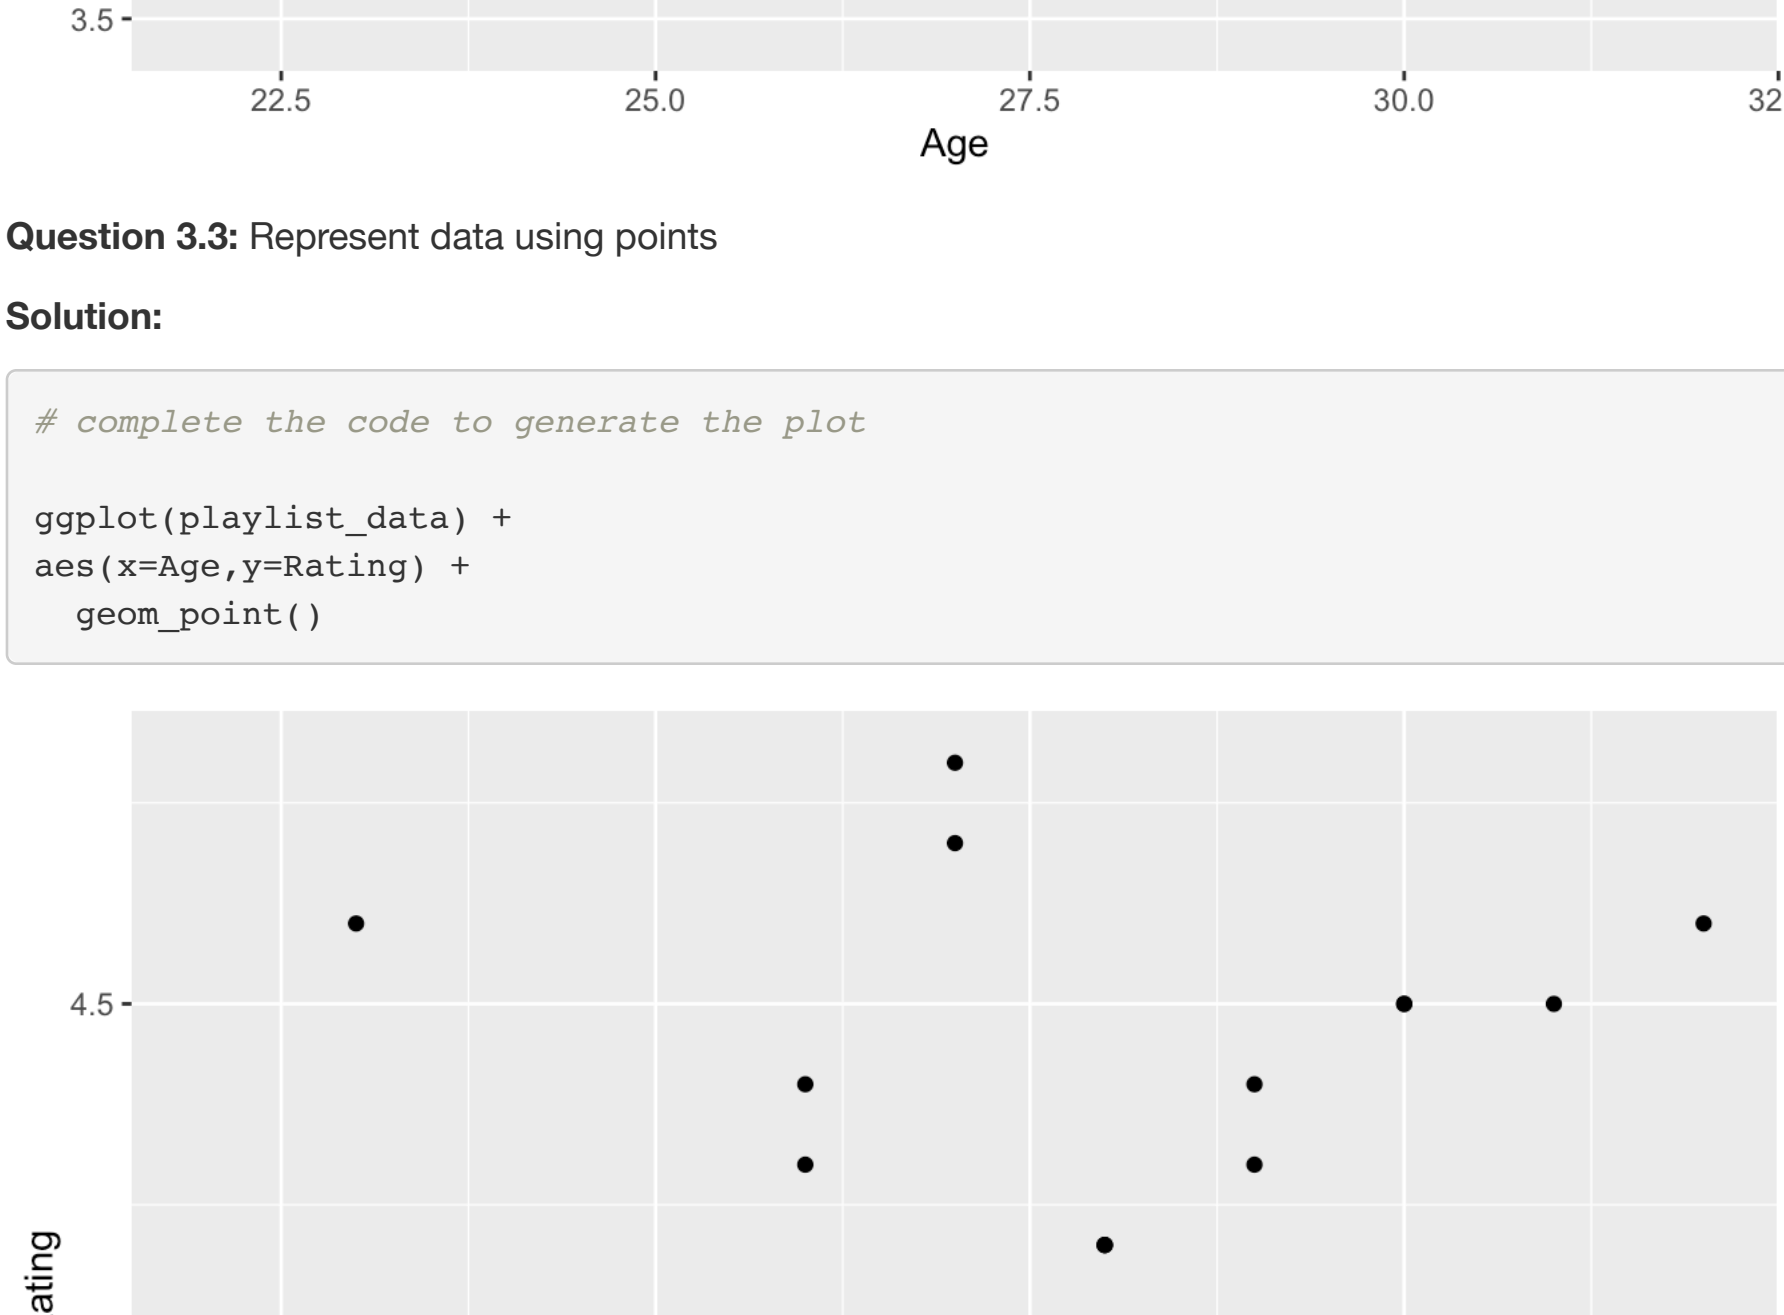
Question 3.2: Label the x-axis as "Age" and the y-axis as "Rating."

Solution:



Question 3.3: Represent data using points

Solution:



Question 3.4: Can you change the points represented by dots/small circles to any other shape of your liking?

Solution:



Question 3.5: Insert a suitable title and briefly provide your insights in the caption

Solution:



Source: tidyverse playlist dataset